

Setup Superset

1. Task requirement

To run the superset with a podman play kube which is generated with the helm chart of the superset.

Superset is a modern data exploration and data visualisation platform. Superset can replace or augment proprietary business intelligence tools for many teams. Superset integrates well with a variety of data sources. Superset provides: A no-code interface for building charts quickly.

2. Environment details

- OS: Ubuntu 20.04

My System configuration

- CPU: Intel Core i5-8350U CPU @ 1.70GHz x 8
- RAM: 8GB (4GB x 2 SODIMM DDR4)
- Storage: 512GB

3. List of tools and technologies

- Podman version 3.4.2
- Helm version 3.12.1

Others

- Vim
- Bash

4. Definition of tools

- **Podman** - It is an open source tool for developing, managing, and running containers on your Linux systems.
- **Helm** - It is a tool that streamlines installing and managing Kubernetes applications.
- **Kubernetes** -
- **Vim** - It is a highly configurable text editor built to make creating and changing any kind of text very efficient
- **Bash** - It is the command line shell that you encounter when you open the terminal on most Unix operating systems, like MacOS and Linux.

5. Command for the setup or configuration

Step 1. Run the following command to install curl.

```
sudo apt install curl
```

- **curl** is a command-line tool that allows you to fetch data from the internet. It's like a web browser for your terminal. You can use it to download files, make web requests, and interact with web services directly from the command line.

Output:

```

maansi@maansi-Standard-PC-Q35-ICH9-2009:~$ sudo apt install curl
[sudo] password for maansi:
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following additional packages will be installed:
  libcurl4
The following NEW packages will be installed:
  curl
The following packages will be upgraded:
  libcurl4
1 upgraded, 1 newly installed, 0 to remove and 234 not upgraded.
Need to get 161 kB/396 kB of archives.
After this operation, 417 kB of additional disk space will be used.
Do you want to continue? [Y/n] y
Get:1 http://in.archive.ubuntu.com/ubuntu focal-updates/main amd64 curl amd64 7.68.0-1ubuntu2.19 [161
kB]
Fetched 161 kB in 2s (93.7 kB/s)
(Reading database ... 182339 files and directories currently installed.)
Preparing to unpack .../libcurl4_7.68.0-1ubuntu2.19_amd64.deb ...
Unpacking libcurl4:amd64 (7.68.0-1ubuntu2.19) over (7.68.0-1ubuntu2.16) ...
Selecting previously unselected package curl.
Preparing to unpack .../curl_7.68.0-1ubuntu2.19_amd64.deb ...
Unpacking curl (7.68.0-1ubuntu2.19) ...
Setting up libcurl4:amd64 (7.68.0-1ubuntu2.19) ...
Setting up curl (7.68.0-1ubuntu2.19) ...
Processing triggers for man-db (2.9.1-1) ...
Processing triggers for libc-bin (2.31-0ubuntu9.9) ...

```

Step 2. Adding Helm GPG Key.

```

$ curl https://baltocdn.com/helm/signing.asc | gpg --dearmor | sudo tee
/usr/share/keyrings/helm.gpg > /dev/null

```

- **curl <https://baltocdn.com/helm/signing.asc:>**
curl : To download the contents of the specified URL, which appears to be the GPG (GNU Privacy Guard) signing key for Helm, a package manager for Kubernetes.
- **gpg --dearmor:** In GPG, "dearmor" means to convert a GPG public or private key from the binary format to a text-based format. This is often done for distribution or storage purposes.
- **sudo tee /usr/share/keyrings/u2`helm.gpg > /dev/null** : This part of the command uses tee to write the output of the previous gpg command to the file /usr/share/keyrings/helm.gpg.
- **> /dev/null** : redirects the standard output (stdout) to the "null" device .

Output:

```

maansi@indianrenters-Latitude-5490:~$ curl https://baltocdn.com/helm/signing.asc | gpg --dearmor | sudo tee /usr/share/keyrings/helm.gpg > /dev/null
[sudo] password for maansi:
% Total    % Received % Xferd Average Speed   Time    Time     Time  Current
   Dload  Upload  Total   Spent    Left     Speed
100 1699 100 1699    0     0  5838      0 --:--:-- --:--:-- --:--:--  5838
maansi@indianrenters-Latitude-5490:~$

```

Step 3. Adding Helm Repository to Package Sources.

```

$ echo "deb [arch=$(dpkg --print-architecture)
signed-by=/usr/share/keyrings/helm.gpg]
https://baltocdn.com/helm/stable/debian/ all main" | sudo tee
/etc/apt/sources.list.d/helm-stable-debian.list

```

- `echo` : display the text/string on the terminal
- line tells the apt package manager to download Helm packages from the <https://baltocdn.com/helm/stable/debian/> repository.
- The `arch=$(dpkg --print-architecture)` : apt package manager to download the correct package for our system architecture.
- The `signed-by=/usr/share/keyrings/helm.gpg` part tells the apt package manager to verify the authenticity of the packages using the GPG key that is stored in the file `/usr/share/keyrings/helm.gpg`.
- Creates a new file called `helm-stable-debian.list` in the directory `/etc/apt/sources.list.d/`.
- It adds the following line to the file: `deb [arch=$(dpkg --print-architecture) signed-by=/usr/share/keyrings/helm.gpg] https://baltocdn.com/helm/stable/debian/ all main`

Output:

```

maansi@indianrenters-Latitude-5490:~$ echo "deb [arch=$(dpkg --print-architecture) signed-by=/usr/share/keyrings/helm.gpg] https://baltocdn.com/helm/stable/debian/ all main" | sudo tee /etc/apt/sources.list.d/helm-stable-debian.list
deb [arch=amd64 signed-by=/usr/share/keyrings/helm.gpg] https://baltocdn.com/helm/stable/debian/ all main
maansi@indianrenters-Latitude-5490:~$

```

Step 4. Updating Packages and Repositories.

```

$ sudo apt update

```

- **sudo:** This stands for "Superuser Do" and is used to execute commands with superuser (administrator) privileges. It allows you to perform actions that require elevated permissions.
- **apt:** This stands for "Advanced Package Tool." It's a command-line tool used to manage software packages on Debian-based systems. You can use it to install, update, remove, and manage software packages.
- **update:** This is a subcommand of **apt** that instructs the package manager to update the package information. The package information is a database of available software packages, their versions, and where to download them from. This database needs to be refreshed periodically to ensure that your system has the latest information about available software.

Output:

```
maansi@indianrenters-Latitude-5490:~$ sudo apt update
Hit:1 https://baltocdn.com/helm/stable/debian all InRelease
Hit:2 http://archive.ubuntu.com/ubuntu focal InRelease
Hit:3 https://packages.cloud.google.com/apt kubernetes-xenial InRelease
Get:4 https://download.opensuse.org/repositories/devel:/kubic:/libcontainers:/stable/xUbuntu_20.04 InRelease [1,642 B]
Fetched 1,642 B in 1s (1,279 B/s)
Reading package lists... Done
Building dependency tree
Reading state information... Done
All packages are up to date.
maansi@indianrenters-Latitude-5490:~$
```

Step 5. Install helm.

```
$ sudo apt install helm
```

Or

```
$ sudo snap install helm --classic
```

- **sudo:** This prefix is used to execute the command with superuser (administrator) privileges. It allows you to install software system-wide.
- **apt:** The package management tool used for managing software packages on Debian-based systems.

- **install**: A command to instruct **apt** to install a specified package.
- **helm**: The name of the package you want to install. In this case, "Helm" is a tool used for managing Kubernetes applications. Helm simplifies the process of deploying and managing complex applications on Kubernetes clusters.

Output:

```
mansi@mansi-Standard-PC-Q35-ICH9-2009:~$ sudo apt install helm
Reading package lists... 0%
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following NEW packages will be installed:
  helm
0 upgraded, 1 newly installed, 0 to remove and 218 not upgraded.
Need to get 16.0 MB of archives.
After this operation, 50.6 MB of additional disk space will be used.
Get:1 https://baltocdn.com/helm/stable/debian all/main amd64 helm amd64 3.12.3-1 [16.0 MB]
Fetched 16.0 MB in 3s (4,813 kB/s)
Selecting previously unselected package helm.
(Reading database ... 179517 files and directories currently installed.)
Preparing to unpack .../helm_3.12.3-1_amd64.deb ...
Unpacking helm (3.12.3-1) ...
Setting up helm (3.12.3-1) ...
Processing triggers for man-db (2.9.1-1) ...
```

Step 6. Verifying the installation and version.

```
$ helm version
```

Output:

```
mansi@mansi-Standard-PC-Q35-ICH9-2009:~$ helm version
version.BuildInfo{Version:"v3.12.3", GitCommit:"3a31588ad33fe3b89af5a2a54ee1d25bfe6eaa5e", GitTreeState:"clean", GoVersion:"go1.20.7"}
mansi@mansi-Standard-PC-Q35-ICH9-2009:~$
```

Step 7. Create directory

```
# mkdir superset-poc
```

- **mkdir**: It's used to create a new directory (folder) in the file system.
- **superset-poc**: This is the name of the directory you want to create.

Output:

Step 8. Go in to directory

```
# cd superset-poc
```

- **cd**: This is a command that stands for "change directory."

- `superset-poc`: This is the name of the directory you want to navigate to.

Output:

```
mansi@mansi-Standard-PC-Q35-ICH9-2009:~$ cd superset-poc
mansi@mansi-Standard-PC-Q35-ICH9-2009:~/superset-poc$
```

Step 9. Run the command to install vim.

```
$ sudo apt install vim
```

Vim is a text editor mainly used for writing and editing text and code. It's efficient, customizable, and works in terminal environments. It's great for programmers, system administrators, and anyone who wants fast and powerful text editing.

Output:

```

maansi@maansi-Standard-PC-Q35-ICH9-2009:~/superset-poc$ sudo apt install vim
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following additional packages will be installed:
  vim-common vim-runtime vim-tiny
Suggested packages:
  ctags vim-doc vim-scripts indent
The following NEW packages will be installed:
  vim vim-runtime
The following packages will be upgraded:
  vim-common vim-tiny
2 upgraded, 2 newly installed, 0 to remove and 232 not upgraded.
Need to get 7,116 kB/7,785 kB of archives.
After this operation, 34.6 MB of additional disk space will be used.
Do you want to continue? [Y/n] y
Get:1 http://in.archive.ubuntu.com/ubuntu focal-updates/main amd64 vim-runtime all 2:8.1.2269-1ubuntu
5.18 [5,875 kB]
Get:2 http://in.archive.ubuntu.com/ubuntu focal-updates/main amd64 vim amd64 2:8.1.2269-1ubuntu5.18 [
1,242 kB]
Fetched 7,116 kB in 3s (2,522 kB/s)
(Reading database ... 182346 files and directories currently installed.)
Preparing to unpack .../vim-tiny_2%3a8.1.2269-1ubuntu5.18_amd64.deb ...
Unpacking vim-tiny (2:8.1.2269-1ubuntu5.18) over (2:8.1.2269-1ubuntu5.11) ...
Preparing to unpack .../vim-common_2%3a8.1.2269-1ubuntu5.18_all.deb ...
Unpacking vim-common (2:8.1.2269-1ubuntu5.18) over (2:8.1.2269-1ubuntu5.11) ...
Selecting previously unselected package vim-runtime.
Preparing to unpack .../vim-runtime_2%3a8.1.2269-1ubuntu5.18_all.deb ...
Adding 'diversion of /usr/share/vim/vim81/doc/help.txt to /usr/share/vim/vim81/doc/help.txt.vim-tiny
by vim-runtime'
Adding 'diversion of /usr/share/vim/vim81/doc/tags to /usr/share/vim/vim81/doc/tags.vim-tiny by vim-r
untime'
Unpacking vim-runtime (2:8.1.2269-1ubuntu5.18) ...
Selecting previously unselected package vim.
Preparing to unpack .../vim_2%3a8.1.2269-1ubuntu5.18_amd64.deb ...
Unpacking vim (2:8.1.2269-1ubuntu5.18) ...
Setting up vim-common (2:8.1.2269-1ubuntu5.18) ...
Setting up vim-runtime (2:8.1.2269-1ubuntu5.18) ...
Setting up vim (2:8.1.2269-1ubuntu5.18) ...
update-alternatives: using /usr/bin/vim.basic to provide /usr/bin/vim (vim) in auto mode
update-alternatives: using /usr/bin/vim.basic to provide /usr/bin/vimdiff (vimdiff) in auto mode
update-alternatives: using /usr/bin/vim.basic to provide /usr/bin/rvim (rvim) in auto mode
update-alternatives: using /usr/bin/vim.basic to provide /usr/bin/rview (rview) in auto mode
update-alternatives: using /usr/bin/vim.basic to provide /usr/bin/vi (vi) in auto mode
update-alternatives: using /usr/bin/vim.basic to provide /usr/bin/view (view) in auto mode
update-alternatives: using /usr/bin/vim.basic to provide /usr/bin/ex (ex) in auto mode
Setting up vim-tiny (2:8.1.2269-1ubuntu5.18) ...
Processing triggers for desktop-file-utils (0.24-1ubuntu3) ...
Processing triggers for mime-support (3.64ubuntu1) ...
Processing triggers for hicolor-icon-theme (0.17-2) ...
Processing triggers for gnome-menus (3.36.0-1ubuntu1) ...
Processing triggers for man-db (2.9.1-1) ...
maansi@maansi-Standard-PC-Q35-ICH9-2009:~/superset-poc$

```

Step 10. Create a yaml file and add details in it .

First I check hosts file and add the domain.

```
sudo vim /etc/hosts
```

And add :

192.168.1.254 fosteringlinux.com and save the file.

Create values-stage.yaml file


```
vim values-stage.yaml
```

- **vim**: This is a command-line text editor that allows you to create, view, and edit text files directly within the terminal.
- **value-stage.yaml**: This is the name of the file you will open.

```
global:
  hosts:
    domain: fosteringlinux.com
postgresql:
  image:
    tag: 14.0
certmanager-issuer:
  email: mansi_01@fosteringlinux.com
```

Note :

"In the values-stage.yaml file, modify the domain and email according to your specific requirements. To save a file in Vim, press Esc and then type :wq! followed by Enter"

"When you copy and paste to another place, the indentation gets messed up, so pay attention to the indentation."

Step 11. Creating Superset Deployment Configuration

```
# helm template --dry-run --debug superset/superset --generate-name
--values values-stage.yaml > superset-kube.yaml
```

- **helm template**: This is the command that tells Helm to generate Kubernetes manifests based on a Helm chart.
- **--dry-run**: This indicates Helm will simulate the installation without actually deploying anything. This is useful for seeing what manifests would be generated without affecting the cluster.
- **--debug**: This enables debug output, providing more detailed information about the Helm template process.

- **superset/superset**: This specifies the Helm chart to use for generating the Kubernetes manifests. The chart name is in the format **repository/chart-name**.
- **--generate-name**: This generates a unique name for the release based on the chart's name. It's used when you want to create a new release without specifying a release name.
- **--values values-stage.yaml**: This points to a values file (**values-stage.yaml**) that provides custom configuration values for the Helm chart.
- **> superset-kube.yaml**: This part of the command uses the output redirection (>) to save the generated Kubernetes manifests to a file named **superset-kube.yaml**.

In summary, this command generates Kubernetes manifests for deploying the Superset application using a specific Helm chart and custom configuration values from the `values-stage.yaml` file. The generated manifests are saved in the `superset-kube.yaml` file. The use of `--dry-run` and `--debug` ensures that we can preview the manifests before actually deploying the application, it helps to verify that the configuration is correct.

Output :

```

Mansi@Mansi-Standard-PC-Q35-ICH9-2009:~/superset-poc$ helm template --dry-run --debug superset/superset --generate-name --values va
lues-stage.yaml > superset-kube.yaml
install.go:200: [debug] Original chart version: ""
Error: repo superset not found
helm.go:84: [debug] repo superset not found
helm.sh/helm/v3/pkg/downloader.pickChartRepositoryConfigByName
helm.sh/helm/v3/pkg/downloader/chart_downloader.go:350
helm.sh/helm/v3/pkg/downloader.(*ChartDownloader).ResolveChartVersion
helm.sh/helm/v3/pkg/downloader/chart_downloader.go:253
helm.sh/helm/v3/pkg/downloader.(*ChartDownloader).DownloadTo
helm.sh/helm/v3/pkg/downloader/chart_downloader.go:90
helm.sh/helm/v3/pkg/action.(*ChartPathOptions).LocateChart
helm.sh/helm/v3/pkg/action/install.go:789
main.runInstall
helm.sh/helm/v3/cmd/helm/install.go:212
main.newTemplateCmd.func2
helm.sh/helm/v3/cmd/helm/template.go:88
github.com/spf13/cobra.(*Command).execute
github.com/spf13/cobra@v1.6.1/command.go:916
github.com/spf13/cobra.(*Command).ExecuteC
github.com/spf13/cobra@v1.6.1/command.go:1044
github.com/spf13/cobra.(*Command).Execute
github.com/spf13/cobra@v1.6.1/command.go:968
main.main
helm.sh/helm/v3/cmd/helm/helm.go:83
runtime.main
runtime/proc.go:250
runtime.goexit
runtime/proc.go:1500

```

Problem 1. There I got the above error .

To solve this I run below commands.

```
$ helm repo add superset https://apache.github.io/superset
```

```
$ helm repo update
```

For this I follow this link:

<https://stackoverflow.com/questions/66706363/where-is-the-superset-helm-chart>

Problem 2.

```
mansi@mansi-Standard-PC-Q35-ICH9-2009:~/superset-poc$ helm template --dry-run --debug superset/superset --generate-name --values values-stage.yaml > superset-kube.yaml
install.go:200: [debug] Original chart version: ""
install.go:217: [debug] CHART PATH: /home/mansi/.cache/helm/repository/superset-0.10.5.tgz
Error: open values-stage.yaml: no such file or directory
helm.go:84: [debug] open values-stage.yaml: no such file or directory
```

When I run this command I get the above error.

To solve this, I change the file name that is incorrect while creating the yaml file. The correct name is **values-stage.yaml** not **value-stage.yaml**.

Then I run the command.

```
# helm template --dry-run --debug superset/superset --generate-name
--values values-stage.yaml > superset-kube.yaml
```

Output:

```
mansi@mansi-Standard-PC-Q35-ICH9-2009:~/superset-poc$ helm template --dry-run --debug superset/superset --generate-name --values values-stage.yaml > superset-kube.yaml
install.go:200: [debug] Original chart version: ""
install.go:217: [debug] CHART PATH: /home/mansi/.cache/helm/repository/superset-0.10.5.tgz
mansi@mansi-Standard-PC-Q35-ICH9-2009:~/superset-poc$
```

Step 12. Edit superset-kube.yaml and create a new final.yaml as per the requirement in podman play kube.

```
cat superset-kube.yaml
```

```
vim final.yaml
```

Add the below data final.yaml file.

```
-----
apiVersion: apps/v1
kind: Deployment
metadata:
  name: kubepostgresql
spec:
```

```
replicas: 1
template:
  metadata:
    name: kubepostgresql
  spec:
    serviceAccountName: default
    securityContext:
      fsGroup: 1001
    hostNetwork: true
    hostIPC: false
    initContainers: []
    containers:
      - name: postgresql
        image: docker.io/bitnami/postgresql:14.6.0-debian-11-r13
        imagePullPolicy: "IfNotPresent"
        securityContext:
          runAsUser: 1001
        env:
          - name: BITNAMI_DEBUG
            value: "false"
          - name: POSTGRESQL_PORT_NUMBER
            value: "5432"
          - name: POSTGRESQL_VOLUME_DIR
            value: "/bitnami/postgresql"
          - name: PGDATA
            value: "/bitnami/postgresql/data"
          - name: POSTGRES_USER
            value: "superset"
          - name: POSTGRES_POSTGRES_PASSWORD
            value: "superset"
          - name: POSTGRES_PASSWORD
            value: "superset"
          - name: POSTGRES_DB
            value: "superset"
          - name: POSTGRESQL_ENABLE_LDAP
            value: "no"
          - name: POSTGRESQL_ENABLE_TLS
            value: "no"
          - name: POSTGRESQL_LOG_HOSTNAME
            value: "false"
          - name: POSTGRESQL_LOG_CONNECTIONS
            value: "false"
          - name: POSTGRESQL_LOG_DISCONNECTIONS
            value: "false"
          - name: POSTGRESQL_PGAUDIT_LOG_CATALOG
            value: "off"
```

```

      # Others
      - name: POSTGRESQL_CLIENT_MIN_MESSAGES
        value: "error"
      - name: POSTGRESQL_SHARED_PRELOAD_LIBRARIES
        value: "pgaudit"
      ports:
      - name: tcp-postgresql
        containerPort: 5432
      resources:
        limits: {}
        requests:
          cpu: 250m
          memory: 256Mi
        volumeMounts:
        - name: data
          mountPath: /bitnami/postgresql
    volumes:
    - name: data
      hostPath:
        path: /home/maansi/superset-poc/postgres/

---

apiVersion: apps/v1
kind: Deployment
metadata:
  name: kuberediis-master
spec:
  replicas: 1
  template:
    spec:
      securityContext:
        fsGroup: 1001
      terminationGracePeriodSeconds: 30
      containers:
      - name: redis
        image: docker.io/bitnami/redis:7.0.10-debian-11-r4
        imagePullPolicy: "IfNotPresent"
        securityContext:
          runAsUser: 1001
        command:
        - /bin/bash
        args:
        - -c
        - /opt/bitnami/scripts/start-scripts/start-master.sh
        env:

```

```
- name: BITNAMI_DEBUG
value: "false"
- name: REDIS_REPLICATION_MODE
value: "master"
- name: ALLOW_EMPTY_PASSWORD
value: "yes"
- name: REDIS_TLS_ENABLED
value: "no"
- name: REDIS_PORT
value: "6379"
ports:
- name: redis
containerPort: 6379
hostPort: 6379
resources:
limits: {}
requests: {}
volumeMounts:
- name: start-scripts
mountPath:
/opt/bitnami/scripts/start-scripts/start-master.sh
subPath: start-master.sh
- name: redis-data
mountPath: /data
- name: config
mountPath: /opt/bitnami/redis/mounted-etc
volumes:
- name: start-scripts
hostPath:
path: /home/maansi/superset-poc/redis/start-master.sh
type: FileOrCreate
defaultMode: 0777
- name: config
hostPath:
path: /home/maansi/superset-poc/redis/redis-conf
type: Directory
- name: redis-data
hostPath:
path: /home/maansi/superset-poc/redis/redis-data

---

apiVersion: apps/v1
kind: Deployment
metadata:
name: kubesuperset-worker
```

```
spec:
  replicas: 1
  template:
    spec:
      securityContext:
        runAsUser: 0
      initContainers:
        - command:
            - /bin/sh
            - -c
            - dockerize -wait "tcp://192.168.122.113:5432" -wait
              "tcp://192.168.122.113:6379" -timeout 120s
          env:
            - name: REDIS_HOST
              value: 192.168.122.113
            - name: REDIS_PORT
              value: "6379"
            - name: DB_HOST
              value: 192.168.122.113
            - name: DB_PORT
              value: "5432"
            - name: DB_USER
              value: "superset"
            - name: DB_PASS
              value: "superset"
            - name: DB_NAME
              value: "superset"
          image: 'apache/superset:dockerize'
          imagePullPolicy: 'IfNotPresent'
          name: wait-for-postgres-redis
      containers:
        - name: superset
          image:
            "apachesuperset.docker.scarf.sh/apache/superset:2.1.0"
          imagePullPolicy: IfNotPresent
          command: ["/bin/sh", "-c", ".
/app/pythonpath/superset_bootstrap.sh; celery
--app=superset.tasks.celery_app:app worker"]
          env:
            - name: "SUPERSET_PORT"
              value: "8088"
            - name: REDIS_HOST
              value: 192.168.122.113
            - name: REDIS_PORT
              value: "6379"
            - name: DB_HOST
```

```

        value: 192.168.122.113
        - name: DB_PORT
        value: "5432"
        - name: DB_USER
        value: "superset"
        - name: DB_PASS
        value: "superset"
        - name: DB_NAME
        value: "superset"
        volumeMounts:
        - name: superset-config
        mountPath: "/app/pythonpath"
        readOnly: true
        livenessProbe:
        exec:
        command:
        - sh
        - -c
        - celery -A superset.tasks.celery_app:app inspect ping
-d celery@$HOSTNAME
        failureThreshold: 3
        initialDelaySeconds: 120
        periodSeconds: 60
        successThreshold: 1
        timeoutSeconds: 60
        resources: {}
volumes:
- name: superset-config
  hostPath:
    path: /home/maansi/superset-poc/superset/

---

apiVersion: apps/v1
kind: Deployment
metadata:
  name: kubesuperset
spec:
  replicas: 1
  template:
    spec:
      securityContext:
        runAsUser: 0
      initContainers:
        - command:
            - /bin/sh

```



```
- -C
- dockerize -wait "tcp://192.168.122.113:5432" -timeout 120s
env:
- name: REDIS_HOST
value: 192.168.122.113
- name: REDIS_PORT
value: "6379"
- name: DB_HOST
value: 192.168.122.113
- name: DB_PORT
value: "5432"
- name: DB_USER
value: "superset"
- name: DB_PASS
value: "superset"
- name: DB_NAME
value: "superset"
image: 'apache/superset:dockerize'
imagePullPolicy: 'IfNotPresent'
name: wait-for-postgres
containers:
- name: superset
image:
"apachesuperset.docker.scarf.sh/apache/superset:2.1.0"
imagePullPolicy: IfNotPresent
command: ["/bin/sh", "-c", ".
/app/pythonpath/superset_bootstrap.sh; /usr/bin/run-server.sh"]
env:
- name: "SUPERSET_PORT"
value: "8088"
- name: REDIS_HOST
value: 192.168.122.113
- name: REDIS_PORT
value: "6379"
- name: DB_HOST
value: 192.168.122.113
- name: DB_PORT
value: "5432"
- name: DB_USER
value: "superset"
- name: DB_PASS
value: "superset"
- name: DB_NAME
value: "superset"
volumeMounts:
- name: superset-config
```

```
    mountPath: "/app/pythonpath"
    readOnly: true
    ports:
      - name: http
        containerPort: 8088
        hostPort: 8088
        protocol: TCP
    resources: {}
  volumes:
    - name: superset-config
      hostPath:
        path: /home/maansi/superset-poc/superset/
```

Note : There i replace the path .

Example:

```
path: /home/harsh/superset/redis/start-master.sh
```

To this

```
path: /home/mansi/superset-poc/redis/start-master.sh
```

Also change the ip with your system ip.

"When you copy and paste to another place, the indentation gets messed up, so pay attention to the indentation."

Then save the file.

Step 13. Check list of file in current directory

```
# ls -lrth
```

The command `ls -lrth` is used to list the contents of a directory:

- `-l`: Detailed (long) format showing permissions, owner, size, and more.
- `-r`: Listing in reverse order.
- `-t`: Sorting by modification time (most recent first).
- `-h`: File sizes in a human-readable format.

Output:

```
mansi@mansi-Standard-PC-Q35-ICH9-2009:~/superset-poc$ ls -lrth
total 44K
-rw-rw-r-- 1 mansi mansi 120 Aug 17 13:34 values-stage.yaml
-rw-rw-r-- 1 mansi mansi 29K Aug 17 13:50 superset-kube.yaml
-rw-rw-r-- 1 mansi mansi 7.5K Aug 17 13:52 final.yaml
mansi@mansi-Standard-PC-Q35-ICH9-2009:~/superset-poc$
```

Step 14. Create a script file and add details in it .

```
vim pre-install-task.sh
```

1. Bash Script: A Bash script is a file containing a sequence of commands written in the Bash scripting language. These commands can include system commands, variable assignments, conditionals, loops, and more. Bash scripts are used to automate tasks, perform system administration, and execute a series of actions.

Add in the below script-----

```
#!/bin/bash

mkdir superset
mkdir redis
mkdir postgres

cat <<- END > superset/superset_bootstrap.sh
    #!/bin/bash
    if [ ! -f ~/bootstrap ]; then echo "Running Superset with uid 0" >
~/bootstrap; fi

END

cat <<- END > superset/superset_config.py
import os
from cachelib.redis import RedisCache

def env(key, default=None):
    return os.getenv(key, default)

MAPBOX_API_KEY = env('MAPBOX_API_KEY', '')
CACHE_CONFIG = {
    'CACHE_TYPE': 'RedisCache',
    'CACHE_DEFAULT_TIMEOUT': 300,
    'CACHE_KEY_PREFIX': 'superset_',
    'CACHE_REDIS_HOST': env('REDIS_HOST'),
    'CACHE_REDIS_PORT': env('REDIS_PORT'),
    'CACHE_REDIS_PASSWORD': env('REDIS_PASSWORD'),
    'CACHE_REDIS_DB': env('REDIS_DB', 1),
}
DATA_CACHE_CONFIG = CACHE_CONFIG

SQLALCHEMY_DATABASE_URI =
f"postgresql+psycopg2://{env('DB_USER')}:{env('DB_PASS')}@{env('DB_HOST'
)}:{env('DB_PORT')}/{env('DB_NAME')}"
SQLALCHEMY_TRACK_MODIFICATIONS = True
SECRET_KEY = env('SECRET_KEY', 'thisISaSECRET_1234')

class CeleryConfig(object):
    CELERY_IMPORTS = ('superset.sql_lab', )
    CELERY_ANNOTATIONS = {'tasks.add': {'rate_limit': '10/s'}}
    BROKER_URL = f"redis://{env('REDIS_HOST')}:{env('REDIS_PORT')}/0"
    CELERY_RESULT_BACKEND =
f"redis://{env('REDIS_HOST')}:{env('REDIS_PORT')}/0"
```

```

CELERY_CONFIG = CeleryConfig
RESULTS_BACKEND = RedisCache(
    host=env('REDIS_HOST'),
    port=env('REDIS_PORT'),
    key_prefix='superset_results'
)
END
cat <<- END > superset/superset_init.sh
#!/bin/sh
set -eu
echo "Upgrading DB schema..."
superset db upgrade
echo "Initializing roles..."
superset init

echo "Creating admin user..."
superset fab create-admin \
    --username admin \
    --firstname Superset \
    --lastname Admin \
    --email admin@superset.com \
    --password admin \
    || true

if [ -f "/app/configs/import_datasources.yaml" ]; then
    echo "Importing database connections.... "
    superset import_datasources -p /app/configs/import_datasources.yaml
fi
END
mkdir redis/redis-conf
cat <<- END > redis/redis-conf/master.conf
    dir /data
    # User-supplied master configuration:
    rename-command FLUSHDB ""
    rename-command FLUSHALL ""
    # End of master configuration
END

cat <<- END > redis/redis-conf/redis.conf
    appendonly yes
    save ""
END

cat <<- END > redis/redis-conf/replica.conf
    dir /data
    rename-command FLUSHDB ""

```

```

        rename-command FLUSHALL ""
END

mkdir -p redis/redis-data
cat <<- END > redis/master-redis.sh
#!/bin/bash
[[ -f \${REDIS_PASSWORD_FILE} ]] && export REDIS_PASSWORD="\${(<
"\${REDIS_PASSWORD_FILE})}"
if [[ -f /opt/bitnami/redis/mounted-etc/master.conf ]];then
cp /opt/bitnami/redis/mounted-etc/master.conf
/opt/bitnami/redis/etc/master.conf
fi
if [[ -f /opt/bitnami/redis/mounted-etc/redis.conf ]];then
cp /opt/bitnami/redis/mounted-etc/redis.conf
/opt/bitnami/redis/etc/redis.conf
fi
ARGS=( "--port" "\${REDIS_PORT}" )
ARGS+=( "--protected-mode" "no" )
ARGS+=( "--include" "/opt/bitnami/redis/etc/redis.conf" )
ARGS+=( "--include" "/opt/bitnami/redis/etc/master.conf" )
exec redis-server "\${ARGS[@]}"
END
chmod +x superset/superset_init.sh
chmod +x superset/superset_config.py
chmod +x superset/superset_bootstrap.sh

```

And Save the file.

Step 15. Check list of file in current directory

```
# ls -lrth
```

Output:

```

mansimansi-Standard-PC-Q35-ICH9-2009:~/superset-poc$ ls -lrth
total 48K
-rw-rw-r-- 1 mansi mansi 120 Aug 17 13:34 values-stage.yaml
-rw-rw-r-- 1 mansi mansi 29K Aug 17 13:50 superset-kube.yaml
-rw-rw-r-- 1 mansi mansi 7.5K Aug 17 13:52 final.yaml
-rw-rw-r-- 1 mansi mansi 3.2K Aug 17 14:05 pre-install-task.sh
mansimansi-Standard-PC-Q35-ICH9-2009:~/superset-poc$ █

```

Step 16. Give permission

```
# chmod +x pre-install-task.sh
```

- **chmod** : change the permissions of a file or directory.
- **+x** : to add the execute permission to the file or directory.

Output:

```
mansl@mansl-Standard-PC-Q35-ICH9-2009:~/superset-poc$ chmod +x pre-install-task.sh
mansl@mansl-Standard-PC-Q35-ICH9-2009:~/superset-poc$
```

Step 17. Run the script

```
sh pre-install-task.sh
```

- **sh**: Command to run a shell interpreter.
- **pre-install-task.sh**: Name of the shell script file to be executed.

Step 18. Step 15. Check list of file in current directory.

```
# ls -lrth
```

Output :

```
maansi@maansi-Standard-PC-Q35-ICH9-2009:~/superset-poc$ ls -lrth
total 64K
-rw-rw-r-- 1 maansi maansi 141 Oct  9 23:50 values-stage.yaml
-rw-rw-r-- 1 maansi maansi 29K Oct  9 23:54 superset-kube.yaml
-rw-rw-r-- 1 maansi maansi 8.2K Oct  9 23:57 final.yaml
-rwxrwxr-x 1 maansi maansi 3.2K Oct  9 23:58 pre-install-task.sh
drwxrwxr-x 2 maansi maansi 4.0K Oct  9 23:59 postgres
drwxrwxr-x 2 maansi maansi 4.0K Oct  9 23:59 superset
drwxrwxr-x 4 maansi maansi 4.0K Oct  9 23:59 redis
maansi@maansi-Standard-PC-Q35-ICH9-2009:~/superset-poc$
```

Step 19. Install podman.

```
sudo apt install podman
```

- **sudo**: Execute with superuser privileges.
- **apt**: Package management tool.
- **install**: Command to install a package.
- **podman**: Name of the package (a containerization tool).

Output:

Problem 3 I got error while installing podman.

```
mansi@mansi-Standard-PC-Q35-ICH9-2009:~/superset-poc$ sudo apt install podman
Reading package lists... Done
Building dependency tree
Reading state information... Done
E: Unable to locate package podman
```

For this i run below commands to solve this problem.

```
$ sudo sh -c "echo 'deb
https://download.opensuse.org/repositories/devel:/kubic:/libcontainers:/
stable/xUbuntu_$(lsb_release -rs)/ /' >
/etc/apt/sources.list.d/devel:kubic:libcontainers:stable.list"
```

```
$ wget
https://download.opensuse.org/repositories/devel:/kubic:/libcontainers:/
stable/xUbuntu_$(lsb_release -rs)/Release.key
```

```
$ sudo apt-key add - < Release.key
```

```
$ sudo apt update
```

```
$ sudo apt install -y podman
```

```
$ podman --version
```

"Note : Adding a repository is required for Ubuntu 20 version and is not needed for the latest version of Ubuntu."

Step 19. Create the pods and containers

```
# podman play kube final.yaml
```

- The command `podman play kube final.yaml` will read the Kubernetes YAML file `final.yaml` and recreate the pods and containers described in the file. create and start pods and containers without having to use the Kubernetes API or kubectl.

Output:

```
maansi@indianrenters-Latitude-5490:~/superset-poc$ podman play kube final.yaml
Trying to pull docker.io/bitnami/postgresql:14.6.0-debian-11-r13...
Getting image source signatures
Copying blob 380df3b45a2b done
Copying blob f8c1c832ce65 done
Copying config ff56fb78c6 done
Writing manifest to image destination
Storing signatures
Trying to pull docker.io/bitnami/redis:7.0.10-debian-11-r4...
Getting image source signatures
Copying blob 66e6a14f5a11 done
Copying config b762b5221e done
Writing manifest to image destination
Storing signatures
Resolved "apache/superset" as an alias (/home/maansi/.cache/containers/short-name-aliases.conf)
Trying to pull docker.io/apache/superset:dockerize...
Getting image source signatures
Copying blob 7264a8db6415 done
Copying blob f8213b5a0d73 done
Copying config 0066af9b14 done
Writing manifest to image destination
Storing signatures
Trying to pull apachesuperset.docker.scarf.sh/apache/superset:2.1.0...
Getting image source signatures
Copying blob 2de5f76669ff done
Copying blob 2d2b01660885 done
Copying blob f1f26f570256 done
Copying blob 2a34e3ad1a24 done
Copying blob 9450a9e89370 done
Copying blob 558c7c82036e done
Copying blob fd6e8aa0d19b done
Copying blob b8c33915aa76 done
Copying blob 4db7a27e56a7 done
Copying blob 1f3e6433dafd done
Copying blob f53268c71308 done
Copying blob c43b42a5c1bd done
Copying blob 31a5b66b66dc done
Copying blob dc53248fbae6 done
Copying blob db0542db1f78 done
Copying config 27d3d51cc6 done
Writing manifest to image destination
Storing signatures
Pod:
13d686ee154a1953c81ae1d94c76669a6ed0c28e584c1eb7bbee35e5150e005c
Container:
ea50f72829059eca40a6caced114f5fae5397d87030058165200300bf3625e452

Pod:
d362247981da2ddd81774c9c9d392b9ed421fb596b3766c75be1824de8db88aa
Container:
7d4a2828f343febb14d9c4cefbd29ebc49f70a5a2812b691b4164967bfaa0df

Pod:
699341b47e5827c4409e833785a45d45da0ae7ed66dffa12093868531d070c9b
Container:
ab047f09c77ecd8f5cd7dbb4b5453eed32eedfe9e6d7b1a5d6d3f04e8d2721ea

Pod:
769d16c0f49af7dd5c594ea9df5dcbdf63af70710d8266235f96e399888e2be4
Container:
732561882f104314b6f4e45e735c6cc335960fad43a4cd7325fff8912c712

maansi@indianrenters-Latitude-5490:~/superset-poc$
```

Step 20. To see running containers.

```
# podman ps
```

- List the currently running containers managed by the Podman tool.
- Provides information about container IDs, names, status, ports, and more.

Output :

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
c6977a321740	k8s.gcr.io/pause:3.5	podman ps	About an hour ago	Up About an hour ago		12b69f139b95-lnfra
34b796198509	docker.io/bitnami/postgresql:14.6.0-debian-11-r13	/opt/bitnami/scr...	About an hour ago	Up About an hour ago		kubepostgresql-pod-0-postgresql
5a8fdd820ab	k8s.gcr.io/pause:3.5		About an hour ago	Up About an hour ago	0.0.0.0:6379->6379/tcp	6cf5b91bcf86-lnfra
1285b091b9a0	k8s.gcr.io/pause:3.5		About an hour ago	Up About an hour ago		cb3402eeefb7-lnfra
eeefeeac48e	apachesuperset.docker.scarf.sh/apache/superset:2.1.0		About an hour ago	Up 15 minutes ago (healthy)		kubesuperset-worker-pod-0-superset
e9553c90fe0d	k8s.gcr.io/pause:3.5		About an hour ago	Up About an hour ago	0.0.0.0:8088->8088/tcp	0d2981fe9acf-lnfra
ea1b1e2bf05e	apachesuperset.docker.scarf.sh/apache/superset:2.1.0		About an hour ago	Up About an hour ago	0.0.0.0:8088->8088/tcp	kubesuperset-pod-0-superset

1) Go to the browser and write :

192.168.122.113:8088/login/

If you encounter error after this follow given below steps.

ls -lrth

```
maansi@indianrenters-Latitude-5490:~/superset-poc$ ls -lrth
total 64K
-rw-rw-r-- 1 maansi maansi 140 Oct 2 03:04 values-stage.yaml
-rw-rw-r-- 1 maansi maansi 29K Oct 2 03:04 superset-kube.yaml
-rwxrwxr-x 2 maansi maansi 4.0K Oct 2 03:08 superset
-rwxrwxr-x 4 maansi maansi 4.0K Oct 2 03:09 redis
-rwxrwxr-x 3 maansi maansi 4.0K Oct 2 03:09 postgres
-rwxrwxr-x 1 maansi maansi 3.2K Oct 2 20:59 pre-install-task.sh
-rw-rw-r-- 1 maansi maansi 8.3K Oct 3 11:05 final.yaml
maansi@indianrenters-Latitude-5490:~/superset-poc$
```

cd superset

```
maansi@indianrenters-Latitude-5490:~/superset-poc$ cd superset/
maansi@indianrenters-Latitude-5490:~/superset-poc/superset$
```

ls -lrth

```
maansi@indianrenters-Latitude-5490:~/superset-poc/superset$ ls -lrth
total 12K
-rwxrwxr-x 1 maansi maansi 96 Oct 2 03:08 superset_bootstrap.sh
-rwxrwxr-x 1 maansi maansi 1.2K Oct 2 03:08 superset_config.py
-rwxrwxr-x 1 maansi maansi 556 Oct 2 03:08 superset_init.sh
maansi@indianrenters-Latitude-5490:~/superset-poc/superset$
```

vim superset_config.py

```
maansi@maansi-Standard-PC-Q35-ICH9-2009: ~/superset-poc/superset
import os
from cachelib.redis import RedisCache

def env(key, default=None):
    return os.getenv(key, default)

MAPBOX_API_KEY = env('MAPBOX_API_KEY', '')
CACHE_CONFIG = {
    'CACHE_TYPE': 'RedisCache',
    'CACHE_DEFAULT_TIMEOUT': 300,
    'CACHE_KEY_PREFIX': 'superset_',
    'CACHE_REDIS_HOST': env('REDIS_HOST'),
    'CACHE_REDIS_PORT': env('REDIS_PORT'),
    'CACHE_REDIS_PASSWORD': env('REDIS_PASSWORD'),
    'CACHE_REDIS_DB': env('REDIS_DB', 1),
}
DATA_CACHE_CONFIG = CACHE_CONFIG

SQLALCHEMY_DATABASE_URI = f"postgresql+psycopg2://{env('DB_USER')}:{env('DB_PASS')}@{env('DB_HOST')}:{env('DB_PORT')}/{env('DB_NAME')}"
SQLALCHEMY_TRACK_MODIFICATIONS = True
SECRET_KEY = env('SECRET_KEY', 'thisISaSECRET_1234')

class CeleryConfig(object):
    CELERY_IMPORTS = ('superset.sql_lab', )
    CELERY_ANNOTATIONS = {'tasks.add': {'rate_limit': '10/s'}}
    BROKER_URL = f"redis://{env('REDIS_HOST')}:{env('REDIS_PORT')}/0"
    CELERY_RESULT_BACKEND = f"redis://{env('REDIS_HOST')}:{env('REDIS_PORT')}/0"

CELERY_CONFIG = CeleryConfig
RESULTS_BACKEND = RedisCache(
    host=env('REDIS_HOST'),
    port=env('REDIS_PORT'),
    key_prefix='superset_results'
)
```

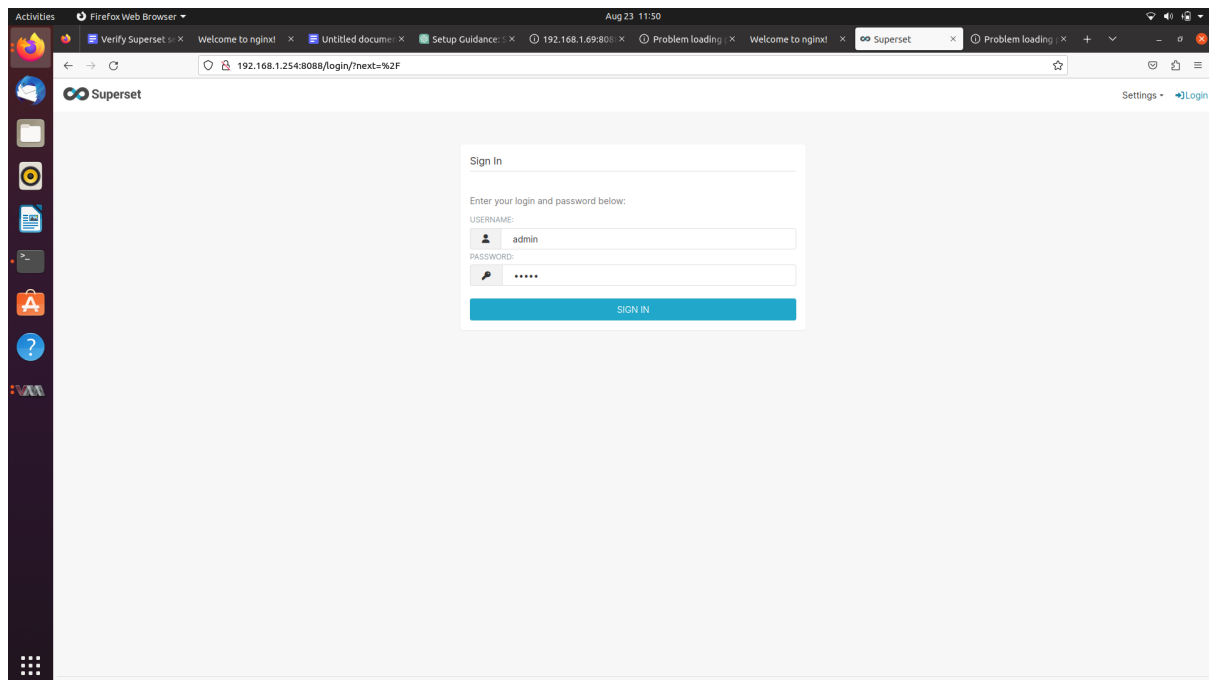
Note :- Correct its indentation and save it.

Test cases list

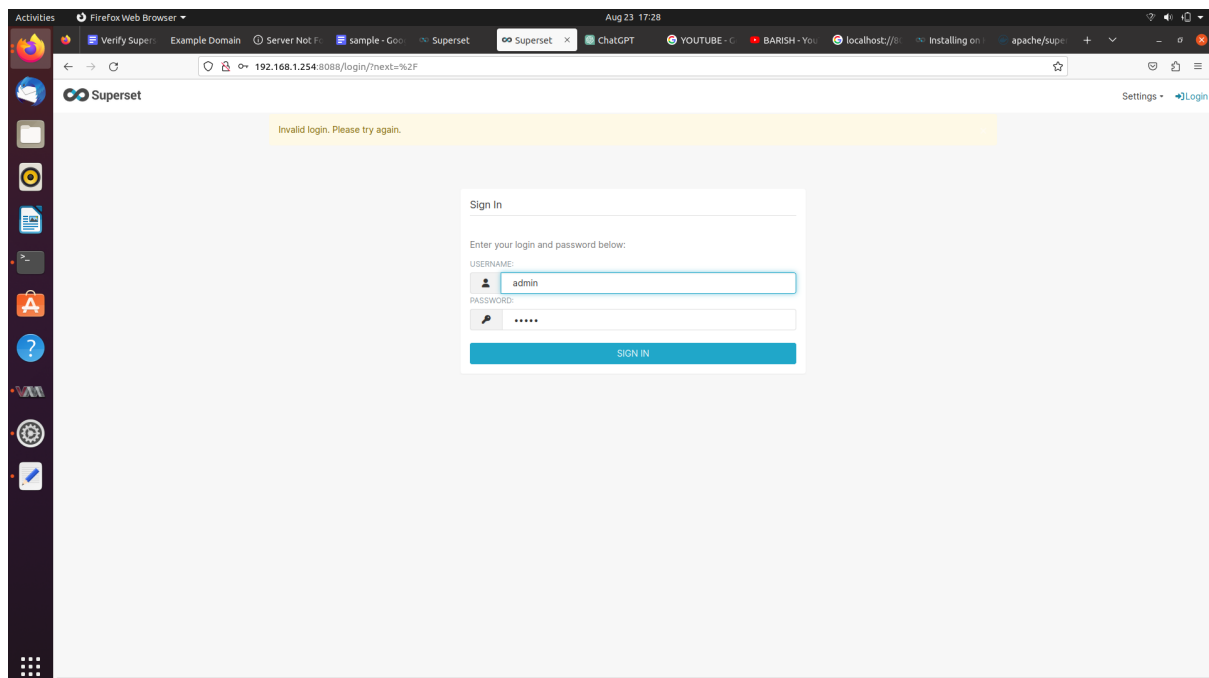
SNO	Component/Tool Name	Test case	Test Count	Test Cases	Expected Result	Test Passed [PASS/FAIL]	Remarks
a	Podman kube	Login to superset	1		Successful login	PASS	
b	Podman kube	Delete the pods and recreate			No data loss	PASS	

2) Go to the browser and write :

192.168.122.113:8088/login/



Problem 4 . When i go the browser and give password and username i got error.



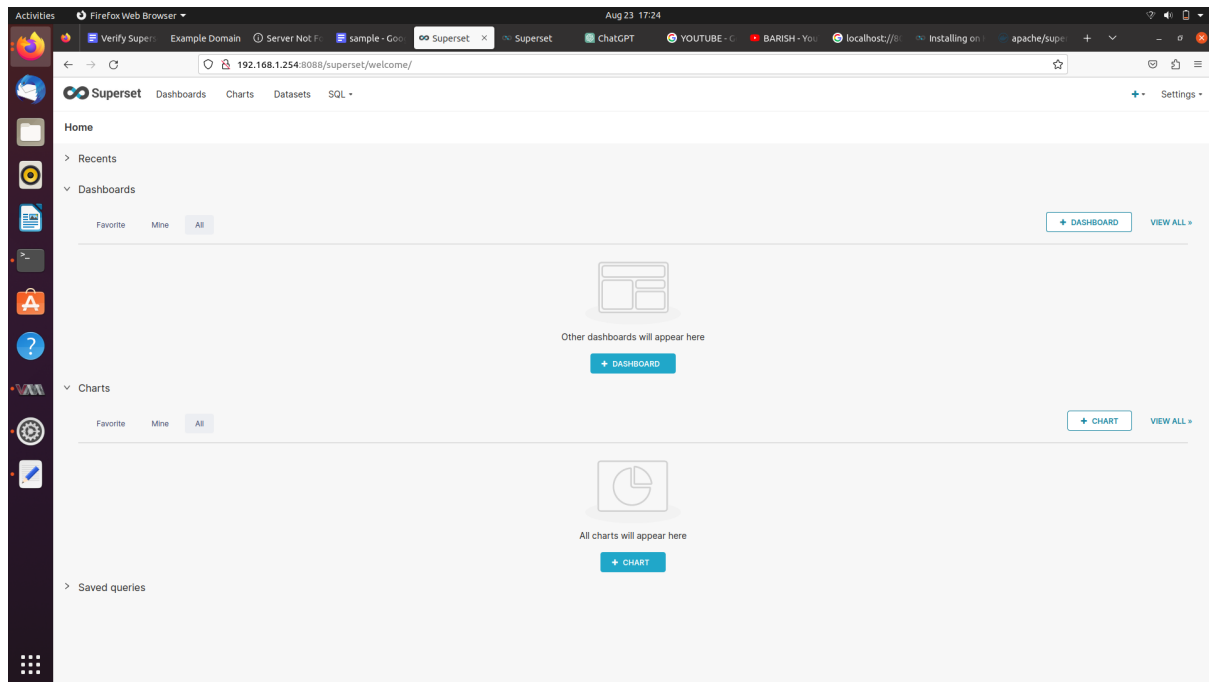
For this error i run below commands.

```
$ podman exec -it kubesuperset-pod-0-superset superset fab create-admin
```

```
$ podman exec -it kubesuperset-pod-0-superset superset db upgrade
```

```
$ podman exec -it kubesuperset-pod-0-superset superset load_examples
```

```
$ podman exec -it kubesuperset-pod-0-superset superset init
```



Reference link

<https://helm.sh/docs/intro/install/>

<https://docs.podman.io/en/v4.2/markdown/podman-play-kube.1.html>